

Experiment-6

Name: -P.Divyashree

Regno: -192472291

Course Code: -MLA0305

Slot-B

Title: - Monte Carlo Prediction and Control for Reinforcement Learning

Aim- To implement Monte Carlo Prediction and Control methods for estimating value functions and finding an optimal policy from sampled episodes.

Procedure

1. Compare the results of Policy Iteration and Value Iteration. Initialize the Reinforcement Learning environment.
2. Define the states, actions, and rewards.
3. Generate an episode by following a selected policy.
4. Store the sequence of:
 - States
 - Actions
 - Rewards
5. Calculate the return G_t from each visited state.
6. Update the estimated value of the state using the observed returns.
7. Repeat the process for multiple episodes.
8. Use the estimated values to improve the policy.
9. Continue generating episodes and updating the policy.
10. Plot the learning performance and cumulative rewards.

Source Code

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.mixture import GaussianMixture
from IPython.display import display, FileLink
from google.colab import files
np.random.seed(42)
print("=" * 80)
print("EXPERIMENT 6")
print("MONTE CARLO PREDICTION AND CONTROL")
print("=" * 80)
N = 500
base = np.column_stack([
    np.random.randint(0, 10, N),
    np.random.randint(0, 2, N),
    np.random.uniform(-5, 10, N),
    np.random.uniform(0.85, 0.99, N)
])
```

```

gmm = GaussianMixture(
    n_components=3,
    random_state=42
)
gmm.fit(base)
generated = gmm.sample(N)[0]
dataset = pd.DataFrame(
    generated,
    columns=[
        "State",
        "Action",
        "Reward",
        "Gamma"
    ]
)
dataset["State"] = dataset["State"].clip(0, 9).round().astype(int)
dataset["Action"] = dataset["Action"].clip(0, 1).round().astype(int)
dataset["Reward"] = dataset["Reward"].clip(-5, 10).round(2)
dataset["Gamma"] = dataset["Gamma"].clip(0.85, 0.99).round(3)
print("\n10 SAMPLE DATASET RECORDS")
display(dataset.head(10))
print("Total Records:", len(dataset))

# =====
# 3. SUMMARY
# =====

display(
    dataset.describe().T[
        ["min", "max", "mean", "std"]
    ].round(3)
)

# =====
# 4. DOWNLOAD
# =====

filename = "Experiment_6_Monte_Carlo_GMM_Dataset.csv"

dataset.to_csv(
    filename,
    index=False
)

print("\nDOWNLOAD DATASET:")

```

```
display(FileLink(filename))

files.download(filename)

# =====
# 5. ENVIRONMENT
# =====

N_STATES = 10
N_ACTIONS = 2

EPISODES = 500
GAMMA = 0.90
ALPHA = 0.10
EPSILON = 0.10

# =====
# 6. EPISODE
# =====

def generate_episode():

    state = 0
    episode = []

    while state != 9:

        action = np.random.randint(
            N_ACTIONS
        )

        if action == 0:
            next_state = min(
                state + 1,
                9
            )
        else:
            next_state = max(
                state - 1,
                0
            )

        reward = (
            10
```

```

        if next_state == 9
            else -1
    )

    episode.append(
        (state, action, reward)
    )

    state = next_state

    return episode

# =====
# 7. MONTE CARLO PREDICTION
# =====

V_mc = np.zeros(N_STATES)

returns_sum = np.zeros(N_STATES)
returns_count = np.zeros(N_STATES)

prediction_history = []

for episode in range(EPISODES):

    ep = generate_episode()

    G = 0
    visited = set()

    for state, action, reward in reversed(ep):

        G = reward + GAMMA * G

        if state not in visited:

            returns_sum[state] += G
            returns_count[state] += 1

            V_mc[state] = (
                returns_sum[state]
                /
                returns_count[state]
            )

```

```

        visited.add(state)

    prediction_history.append(
        np.mean(V_mc)
    )

# =====
# 8. MONTE CARLO CONTROL
# =====

Q = np.zeros(
    (N_STATES, N_ACTIONS)
)

control_rewards = []

for episode in range(EPISODES):

    state = 0
    ep = []

    while state != 9:

        if np.random.rand() < EPSILON:
            action = np.random.randint(
                N_ACTIONS
            )
        else:
            action = np.argmax(
                Q[state]
            )

        if action == 0:
            next_state = min(
                state + 1,
                9
            )
        else:
            next_state = max(
                state - 1,
                0
            )

        reward = (
            10

```

```

        if next_state == 9
        else -1
    )

    ep.append(
        (state, action, reward)
    )

    state = next_state

G = 0

for s, a, r in reversed(ep):

    G = r + GAMMA * G

    Q[s, a] += (
        ALPHA *
        (G - Q[s, a])
    )

    control_rewards.append(
        sum(x[2] for x in ep)
    )

# =====
# 9. RESULT TABLE
# =====

state_values = pd.DataFrame({
    "State": range(N_STATES),
    "MC_Value": V_mc,
    "Best_Control_Value": np.max(Q, axis=1),
    "Best_Action": np.argmax(Q, axis=1)
}).round(3)

print("\nMONTE CARLO STATE VALUES")

display(state_values)

# =====
# 10. SUMMARY
# =====

summary = pd.DataFrame({

```

```

        "Metric": [
            "Dataset Records",
            "Episodes",
            "States",
            "Actions",
            "Gamma",
            "Alpha",
            "Epsilon",
            "Average MC Value",
            "Maximum MC Value",
            "Average Control Reward",
            "Maximum Control Reward"
        ],

        "Result": [
            len(dataset),
            EPISODES,
            N_STATES,
            N_ACTIONS,
            GAMMA,
            ALPHA,
            EPSILON,
            round(np.mean(V_mc), 3),
            round(np.max(V_mc), 3),
            round(np.mean(control_rewards), 3),
            max(control_rewards)
        ]
    })

print("\nSUMMARY RESULTS")

display(summary)

# =====
# 11. VISUALIZATION 1
# DATASET REWARD
# =====

plt.figure(figsize=(10, 6))

plt.hist(
    dataset["Reward"],
    bins=25,
    color="#f06595",
    edgecolor="black"

```

```

)

plt.title(
    "GMM Synthetic Reward Distribution",
    fontsize=16,
    fontweight="bold"
)

plt.xlabel("Reward")
plt.ylabel("Frequency")

plt.show()

# =====
# 12. VISUALIZATION 2
# MC VALUE
# =====

plt.figure(figsize=(11, 6))

plt.bar(
    range(N_STATES),
    V_mc,
    color="#7950f2",
    edgecolor="black"
)

plt.title(
    "Monte Carlo Estimated State Values",
    fontsize=16,
    fontweight="bold"
)

plt.xlabel("State")
plt.ylabel("Value")

plt.xticks(range(N_STATES))

plt.show()

# =====
# 13. VISUALIZATION 3
# PREDICTION LEARNING
# =====

```



```

plt.figure(figsize=(11, 6))

plt.plot(
    prediction_history,
    color="#20c997",
    linewidth=3
)

plt.title(
    "Monte Carlo Prediction Learning Progress",
    fontsize=16,
    fontweight="bold"
)

plt.xlabel("Episode")
plt.ylabel("Average State Value")

plt.grid(alpha=0.25)

plt.show()

# =====
# 14. VISUALIZATION 4
# CONTROL REWARDS
# =====

plt.figure(figsize=(11, 6))

plt.plot(
    control_rewards,
    color="#ff922b",
    alpha=0.6
)

moving_average = pd.Series(
    control_rewards
).rolling(25).mean()

plt.plot(
    moving_average,
    color="#d9480f",
    linewidth=3,
    label="25-Episode Moving Average"
)

```

```
plt.title(
    "Monte Carlo Control Learning Performance",
    fontsize=16,
    fontweight="bold"
)

plt.xlabel("Episode")
plt.ylabel("Reward")

plt.legend()
plt.grid(alpha=0.25)

plt.show()
plt.figure(figsize=(10, 6))

plt.bar(
    [
        "Average MC Prediction",
        "Average MC Control Reward"
    ],
    [
        np.mean(V_mc),
        np.mean(control_rewards)
    ],
    color=["#4dabf7", "#ff6b6b"],
    edgecolor="black"
)

plt.title(
    "Monte Carlo Prediction vs Control",
    fontsize=16,
    fontweight="bold"
)
plt.ylabel("Performance")

plt.xticks(rotation=10)

plt.show()

print("\nExperiment 6 completed successfully.")
```

10 SAMPLE DATASET RECORDS

	State	Action	Reward	Gamma
0	8	0	-5.00	0.954
1	6	1	0.47	0.850
2	7	0	3.74	0.926
3	5	1	1.80	0.895
4	5	0	-0.70	0.921
5	8	0	2.08	0.908
6	9	1	0.37	0.885
7	5	1	-2.69	0.850
8	8	1	2.57	0.918
9	7	0	-0.45	0.898

FIG: - 1 SAMPLE DATA

min max mean std

State	0.00	9.00	4.302	2.746
Action	0.00	1.00	0.462	0.499
Reward	-5.00	10.00	2.611	4.346
Gamma	0.85	0.99	0.923	0.038

FIG: - 2 MEAN

0	Dataset Records	500.000
1	Episodes	500.000
2	States	10.000
3	Actions	2.000
4	Gamma	0.900
5	Alpha	0.100
6	Epsilon	0.100
7	Average MC Value	0.342
8	Maximum MC Value	10.000
9	Average Control Reward	0.652
10	Maximum Control Reward	2.000

FIG: - 3 SUMMARY RESULTS




FIG: - 4 DATA SET REWARD DISTRIBUTION




FIG: - 5 VISUALIZATION OF ALGORITHM COMPARISON




FIG: - 6 MONTY CARLO PREDICTION




FIG: - 7 MC PERFORMANCE




FIG: - 8 MCPREDICTION VS CONTROL

Result

Monte Carlo Prediction and Control were successfully implemented. The agent learned from complete episodes and gradually improved its policy based on the estimated returns.